Lecture 2: Investigating objects

Managing and Manipulating Data Using R

What we will do today

Libraries we will use today

"Load" the package we will use today (output omitted)

```
library(tidyverse)
```

If package not yet installed, then must install before you load. Install in "console" rather than .Rmd file

Generic syntax: `install.packages("package_name")`

Install "tidyverse": `install.packages("tidyverse")`

Note: when we load package, name of package is not in quotes; but when we install package, name of package is in quotes:

```
install.packages("tidyverse")
```

```
library(tidyverse)
```

1 R Markdown

What is R Markdown

R Markdown documents embed R code, output associated with R code, and text into one document

An R Markdown document is a " 'Living' document that updates every time you compile ["knit"] it"

R Markdown documents have the extension .Rmd

> can think of them as text files with the extension .Rmd rather than .txt

At top of .Rmd file you specify the "output" style, which dictates what kind of formatted document will be created

> e.g., `html_document` or `pdf_document`

When you compile ["knit"] a .Rmd file, the resulting formatted document can be an HTML document, a PDF document, an MS Word document, or many other types

This slide borrows from Darin Christensen

How people use R Markdown

RMarkdown creates many types of static and dynamic/interactive documents

Example of static policy report

Example of dynamic/interactive presentation

How I use R Markdown

journal manuscripts; reports; presentations; for taking notes when I am learning new methods or reading an empirical paper

How we will be using R Markdown files in this class:

homework you submit will be .Rmd files, with "output" style will be `html_document` or `pdf_document`

lectures we write are .Rmd files, where we the output style will be `beamer_presentation` or `html_document`

`beamer_presentation` is essentially a PDF document, where each page is a slide

Creating RMarkdown documents

Do this with a partner

Approach for creating a RMarkdown document.

1. Point-and-click from within RStudio

    Click on *File* » *New File* » *R Markdown* » *Document* » choose *HTML* » click *OK*

    - optional: add title (this is not the file name, just what appears at top of document)
    - optional: add author name

    save the .Rmd file; *File* » *Save As*

    - any file name
    - recommend you save it in same folder you saved this lecture

    "Knit" the entire .Rmd file

    - point-and-click OR shortcut: **Cmd/Ctrl + Shift + k**

Components of a .Rmd file

An RMarkdown (.Rmd) file consists of several parts

1. **YAML header**

   YAML stands for "yet another markup language"

   controls settings that apply to the whole document (e.g., "output" should be html_document or pdf_document, whether to include table of contents, etc.)

   YAML header goes at very top of document

   starts with a line of three horizontal dashes `---` ; ends with a line of three horizontal dashes `---`

2. **Text** in body of .Rmd file

   e.g., headings; description of results, etc.

3. **R code chunks** in body of .Rmd file

```r
a <- c(2,4,6)
a
a-1
```

4. **R output** associated with code chunks

```
#> [1] 2 4 6
#> [1] 1 3 5
```

Comment: Running R code chunks vs. "knit" entire .Rmd file

Two ways to execute R commands in .Rmd file:

1. "Knit" entire .Rmd file

   shortcut: **Cmd/Ctrl + Shift + k**

2. "Run" code chunk or selected lines within code chunk

   Run selected line(s): **Cmd/Ctrl + Enter**

   Run current chunk: **Cmd/Ctrl + Shift + Enter**

Comment on default settings for RStudio:

When you knit entire .Rmd file, "objects" created within .Rmd file will not be available after file comples

When you run code chunk (or selected lines in chunk), objects created by lines you run will be in your "environment" until you remove them or quit R session

Output types of .Rmd file

Common/important output types:

html_document: R Markdown originally designed to create HTML documents

- Most features/code in .Rmd files were written for html_document
- many of these features are available in other output types
- When learning R Markdown, best to start by learning html_document

pdf_document: Requires installation of `tinytex` R package or LaTeX (MiKTeX/MacTeX)

- How it works:
  - You write .Rmd code;
  - When you compile, this .Rmd code is transformed into LaTeX code
  - LaTeX "engine" creates the formatted .pdf file
- Can include some of the same features available for *html_document*
- Can insert LaTeX commands in .Rmd file with *pdf_document* output

beamer_presentation: Requires installation of LaTeX

- "beamer" is the name for presentations written in LaTeX
- essentially creates PDF of presentation slides
- Lectures for this class created with *beamer_presentation* output
- note: YAML header includes `beamer_header.tex` file, which creates some formatting rules and additional commands

Learning more about R Markdown

Resources

Cheat sheets and quick reference:
Cheat Sheet
Quick Reference [I prefer the quick reference]

Chapters/books
Chapter 27 of "R for Data Science" book
R Markdown: The Definative Guide book [I prefer this book]

How you will learn R Markdown

Lectures written as .Rmd file
During class run "code chunks" and try to "knit" entire .Rmd file

I'll assign **small** amount of reading on R Markdown
prior to next week:
spend 15 minutes familiarizing yourself with Quick Reference
Read section 3.1 of R Markdown: The Definative Guide, about creating html_document

Homework must be written in .Rmd file
you submit .Rmd file AND output of compiled file
for next week, you will submit homework as html_document output

Directory structure for this class

[XIN - PLEASE REVISE/EDIT THIS SLIDE] In order to be able to "knit" entire lectures [rather than just run specific code chunks] make sure that you have the following directory structure:

rclass

    lectures

        lecture1

        ...

        lecture10

        beamer_header.tex

What is beamer_header.tex?

A text file that contains LaTeXcode

This code creates formatting rules that are applied to all lecture slides

If you go YAML header you will see:

```
includes:
  in_header: ../beamer_header.tex
```

This runs beamer_header.tex; assumes that beamer_header.tex is located one level up from your current directory

If you don't have beamer_header.tex saved to appropriate place, you can download it here [XIN ADD LINK TO FILE]

Note: we may revise beamer_header.tex as we work out formatting bugs

2 Using functions

What are functions

Functions are pre-written bits of code that accomplish some task.

Functions generally follow three sequential steps:

1. take in an **input** object(s)
2. **process** the input.
3. **return** (A) a new object or (B) a visualizatoin (e.g., plot)

For example, `sum()` function calcualtes sum of elements in a vector

1. **input**. takes in a vector of elements (numeric or logical)
2. **processing**. Calculates the sum of elements
3. **return**. Returns numeric vector of length=1; value is sum of input vector

```
sum(c(1,2,3))
#> [1] 6
typeof(sum(c(1,2,3))) # type of object created by sum()
#> [1] "double"
length(sum(c(1,2,3))) # length of object created by sum()
#> [1] 1

#sum(c(TRUE,TRUE,FALSE))
#typeof(sum(c(TRUE,TRUE,FALSE))); length(sum(c(TRUE,TRUE,FALSE)))
```

Function syntax

Components of a function

function name (e.g., `sum()` , `length()` , `seq()`)

function arguments

Inputs that the function takes, which determine what function does

can be vectors, data frames, logical statements, etc.

In "function call" you specify values to assign to these function arguments

e.g., `sum(c(1,2,3))`

Separate arguments with a comma `,`

e.g., `seq(10,15)` Example: the sequence function, `seq()`

```
seq(10,15)
#> [1] 10 11 12 13 14 15
```

Function syntax: More on function arguments

Usually, function arguments have names

e.g., the `seq()` function includes the arguments `from`, `to`, `by`

when you call the function, you need to assign values to these arguments; but you usually don't have to specify the name of the argument

```
seq(from=10, to=20, by=2)
#> [1] 10 12 14 16 18 20
seq(10,20,2)
#> [1] 10 12 14 16 18 20
```

Many function arguments have "default values", set by whoever wrote function

if you don't specify a value for that argument, the default value is inserted

e.g., partial list of default values for `seq()`: `seq(from=1, to=1, by=1)`

```
seq()
#> [1] 1
seq(to=10)
#> [1] 1 2 3 4 5 6 7 8 9 10
seq(10) # R assigned value of 10 to "to" rather than "from" or "by"
#> [1] 1 2 3 4 5 6 7 8 9 10
```

Function arguments, the `na.rm` argument

When R performs calculation and an input has value `NA`, output value is `NA`

```
5+4+NA
#> [1] NA
```

R functions that perform calculations often have argument named `na.rm`

`na.rm` argument asks whether to remove `NA` values prior to calculation

For most functions, default value is `na.rm = FALSE`

This means "do not remove `NAs`" prior to calculation

e.g., default values for `sum()` function: `sum(..., na.rm = FALSE)`

```
sum(c(1,2,3,NA), na.rm = FALSE) # default value
#> [1] NA
sum(c(1,2,3,NA))
#> [1] NA
```

if you specify, `na.rm = TRUE`, `NA` values removed prior to calculation

```
sum(c(1,2,3,NA), na.rm = TRUE)
#> [1] 6
```

Help files for functions

To see help file on a function, type `?function_name` without parentheses

```
?sum
?seq
```

Contents of help files

>**Description**. What the function does
>
>**Usage**. Syntax, including default values for arguments
>
>**Arguments**. Description of function arguments
>
>**Details**. Details and idiosyncracies of about how the function works.
>
>**Value**. What (object) the function "returns"
>
>> e.g., `sum()` returns vector of length 1 whose value is sum of input vector
>
>**References**. Additional reading
>
>**See Also**. Related functions
>
>**Examples**. Examples of function in action
>
>Bottom of help file identifies the package the function comes from

Practice!

>when you encounter new function, spend two minutes reading help file
>
>over time, help files will feel less cryptic and will start to feel helpful

Function arguments, the dot-dot-dot (...) argument

On help file for many functions, you will see an argument called ... , referred to as the "dot-dot-dot" argument

```
?sum
?seq
```

"Dot-dot-dot" arguments have several uses. What you should know for now:

... refers to arguments that are "un-named"; but user can specify values

e.g., default syntax for sum() : sum(..., na.rm = FALSE)

argument na.rm is "named" (name is na.rm); argument ... un-named

... used to allow a function to take an arbitrary number of arguments:

```
#Here, sum function takes 1 un-named argument, specifically c(10,5,NA)
sum(c(10,5,NA),na.rm=TRUE)
#> [1] 15

#Here the sum function takes 3 un-named arguments
sum(10,5,NA,na.rm=TRUE)
#> [1] 15

#Here the sum function takes 5 un-named arguments
sum(10,5,10,20,NA,na.rm=TRUE)
#> [1] 45
```

3 Investigating objects [base R]

Load .Rdata data frames we will use today

Data on off-campus recruiting events by public universities

Data frame object `df_event`

One observation per university, recruiting event

Data frame object `df_school`

One observation per high school (visited and non-visited)

```
rm(list = ls()) # remove all objects in current environment

getwd()
#> [1] "C:/Users/ozanj/Documents/rclass/lectures/lecture2"
#load dataset with one obs per recruiting event
load(url("https://github.com/ozanj/rclass/raw/master/data/recruiting/recruit_eve
#load("../../data/recruiting/recruit_event_somevars.Rdata")

#load dataset with one obs per high school
load(url("https://github.com/ozanj/rclass/raw/master/data/recruiting/recruit_sch
#load("../../data/recruiting/recruit_school_somevars.Rdata")
```

Listing objects

Files in your working directory

`list.files()` function lists files in your current working directory

if you run this code from .Rmd file, working directory is location .Rmd file is stored

```
getwd() # what is your current working directory
#> [1] "C:/Users/ozanj/Documents/rclass/lectures/lecture2"
list.files()
#>  [1] "fp1.JPG"                        "fp2.JPG"
#>  [3] "lecture2.1_ucla.pdf"            "lecture2.1_ucla.Rmd"
#>  [5] "lecture2.1_ucla.tex"            "lecture2.pdf"
#>  [7] "lecture2.Rmd"                   "lecture2.tex"
#>  [9] "lecture2_test.Rmd"             "lecture2_test.tex"
#> [11] "lecture2_ucla.pdf"             "lecture2_ucla.Rmd"
#> [13] "lecture2_ucla.tex"             "problemset2.html"
#> [15] "problemset2.Rmd"               "problemset2_solutions.html"
#> [17] "problemset2_solutions.html.zip" "problemset2_solutions.Rmd"
#> [19] "sample.html"                   "sample.Rmd"
#> [21] "sample_simple_rmarkdown.txt"   "text"
#> [23] "transform-logical.png"
```

Objects currently open in your R session

Listing objects currently open in your R session

`ls()` function lists objects currently open in R

```r
x <- "hello!"
ls() # Objects open in R
#> [1] "df_event"  "df_school" "x"
```

Removing objects currently open in your R session

`rm()` function removes specified objects open in R

```r
rm(x)
ls()
#> [1] "df_event"  "df_school"
```

Command to remove all objects open in R (I don't run it)

```r
rm(list = ls())
```

Describing objects, focus on **data frames**

type and **length** of a data frame object

Recall that a data frame is an object where **type** is a list

Length of an object is the number of elements

When object is a data frame, number of elements = number of variables

```
typeof(df_event)
#> [1] "list"
length(df_event) # = num elements = num columns
#> [1] 33
```

Number of **columns** and **rows** of data frame object

number of columns = number of elements = number of variables

number of rows = number of observations

```
ncol(df_event) # num columns = num variables
#> [1] 33
nrow(df_event) # num rows = num observations
#> [1] 18680
dim(df_event) # shows number rows by columns
#> [1] 18680    33
```

`str()` provides compact information on structure any object (output omitted)

```
str(df_event)
```

3.1 Variables names

Variable names

`names()` function lists names of elements in an object
```
?names
```

Recall that a data frame is an object where **type** is a **list** and each **element** is **named**

When object is a data frame:

  each element is a variable

  each element name is a variable name

```
names(df_event)
#>  [1] "instnm"              "univ_id"             "instst"
#>  [4] "pid"                 "event_date"          "event_type"
#>  [7] "zip"                 "school_id"           "ipeds_id"
#> [10] "event_state"         "event_inst"          "med_inc"
#> [13] "pop_total"           "pct_white_zip"       "pct_black_zip"
#> [16] "pct_asian_zip"       "pct_hispanic_zip"    "pct_amerindian_zip"
#> [19] "pct_nativehawaii_zip" "pct_tworaces_zip"   "pct_otherrace_zip"
#> [22] "fr_lunch"            "titlei_status_pub"   "total_12"
#> [25] "school_type_pri"     "school_type_pub"     "g12offered"
#> [28] "g12"                 "total_students_pub"  "total_students_pri"
#> [31] "event_name"          "event_location_name" "event_datetime_start"
```

Variable names

Refer to specific named elements of an object using this syntax:

```
object_name$element_name
```

When object is data frame, refer to specific variables using this syntax:

```
data_frame_name$varname
```

This approach to isolating variables very useful for investigating data

```r
typeof(df_event$instnm)
#> [1] "character"
typeof(df_event$med_inc)
#> [1] "double"
```

Variable names

Data frames are lists with following criteria:

each element of list is a vector; each element of list is a variable

length of data frame = number of variables

```
length(df_event)
#> [1] 33
nrow(df_event)
#> [1] 18680
#str(df_event)
```

each element of the list (i.e., variable) has the same length

Length of each variable is equal to number of observations in data frame

```
typeof(df_event$event_state)
#> [1] "character"
length(df_event$event_state)
#> [1] 18680
str(df_event$event_state)
#>  chr [1:18680] "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" ...

typeof(df_event$med_inc)
#> [1] "double"
length(df_event$med_inc)
#> [1] 18680
str(df_event$med_inc)
#>  num [1:18680] 71714 89122 70137 70137 71024 ...
```

Variable names

The object `df_school` has one obs per high school

variable `visits_by_100751` shows number of visits by University of Alabama to each high school

like all variables in a data frame, the var `visits_by_100751` is just a vector

```
typeof(df_school$visits_by_100751)
#> [1] "integer"
length(df_school$visits_by_100751) # num elements in vector = num obs
#> [1] 21301
str(df_school$visits_by_100751)
#>  int [1:21301] 0 0 0 0 0 0 0 0 0 0 ...
sum(df_school$visits_by_100751) # sum of values of var across all obs
#> [1] 3338
```

We perform calculations on a variable like we would on any vector of same type

```
v <- c(2,4,6)
typeof(v)
#> [1] "double"
length(v)
#> [1] 3
sum(v)
#> [1] 12
```

3.2 View and print data

Viewing and printing data frames

Three ways to view/print a data frame object

1. Simply type the object name (output omitted)

   number of observations and rows printed depend on YAML header settings and on attributes (discussed next week) of the object

```
df_event
```

2. Use the `View()` function to view data in a browser

```
View(df_event)
```

3. `head()` to show the first *n* rows

```
#?head
head(df_event, n=5)
```

Viewing and printing data frames

`obj_name[<rows>,<cols>]` to print specific rows and columns of data frame

  particularly powerful when combined with sequences (e.g., `1:10`)

Examples:

  Print first five rows, all vars

```
df_event[1:5, ]
```

  Print first five rows and first three columns

```
df_event[1:5, 1:3]
```

  Print first three columns of the 100th observation

```
df_event[100, 1:3]
```

  Print the 50th observation, all variables

```
df_event[50,]
```

Viewing and printing data

type `obj_name$var_name` to print specific elements (i.e., vars) in data frame

`df_event$zip`

recall that these elements are vectors, with length = number of obs

```
typeof(df_event$zip)
#> [1] "character"
length(df_event$zip)
#> [1] 18680
```

`obj_name$var_name` syntax can be combined with sequences

vectors don't have "rows" or "columns"; they just have elements

so use sequence to identify which elements you want to print

```
df_event$event_state[1:10]
#>  [1] "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA"
df_event$event_type[6:10]
#> [1] "private hs" "private hs" "public hs"  "private hs" "public hs"
```

Can also print multiple variables using `combine()` function

```
c(df_event$event_state[1:5],df_event$event_type[1:5])
#>  [1] "MA"         "MA"         "MA"         "MA"         "MA"
#>  [6] "public hs"  "public hs"  "public hs"  "public hs"  "public hs"
```

Exercise

Create a printing exercise using the df_school data frame

1. Use `obj_name[<rows>,<cols>]` to print the first 5 rows and 3 columns of data frame
2. Use head() to print first 4 observations
3. Use `obj_name$var_name[1:10]` to print the first 10 observations of a variable
4. Use combine() to print the first 3 observations of variables "school_type" & "name"

Solution

1. Use `obj_name[<rows>,<cols>]` to print the first 5 rows and 3 columns of data frame

```
df_school[1:5,1:3]
#> # A tibble: 5 x 3
#>   state_code school_type ncessch
#>   <chr>      <chr>       <chr>
#> 1 AK         public      020000100208
#> 2 AK         public      020000100211
#> 3 AK         public      020000100212
#> 4 AK         public      020000100213
#> 5 AK         public      020000300216
```

Solution

2. Use head() to print first 4 observations

```
head(df_school, n=4)
#> # A tibble: 4 x 26
#>   state_code school_type ncessch name  address city  zip_code pct_white
#>   <chr>      <chr>       <chr>   <chr> <chr>   <chr> <chr>        <dbl>
#> 1 AK         public      020000~ Beth~ 1006 R~ Beth~ 99559         11.8
#> 2 AK         public      020000~ Ayag~ 106 Vi~ Kong~ 99559           0
#> 3 AK         public      020000~ Kwig~ 108 Vi~ Kwig~ 99622           0
#> 4 AK         public      020000~ Nels~ 118 Vi~ Toks~ 99637           0
#> # ... with 18 more variables: pct_black <dbl>, pct_hispanic <dbl>,
#> #   pct_asian <dbl>, pct_amerindian <dbl>, pct_other <dbl>,
#> #   num_fr_lunch <dbl>, total_students <dbl>, num_took_math <dbl>,
#> #   num_prof_math <dbl>, num_took_rla <dbl>, num_prof_rla <dbl>,
#> #   avgmedian_inc_2564 <dbl>, visits_by_110635 <int>,
#> #   visits_by_126614 <int>, visits_by_100751 <int>, inst_110635 <chr>,
#> #   inst_126614 <chr>, inst_100751 <chr>
```

Solution

3. Use `obj_name$var_name[1:10]` to print the first 10 observations of a variable

```
df_school$name[1:10]
#> [1] "Bethel Regional High School" "Ayagina'ar Elitnaurvik"
#> [3] "Kwigillingok School"         "Nelson Island Area School"
#> [5] "Alakanuk School"             "Emmonak School"
#> [7] "Hooper Bay School"           "Ignatius Beans School"
#> [9] "Pilot Station School"        "Kotlik School"
```

Solution

4. Use combine() to print the first 3 observations of variables "school_type" & "name"

```
c(df_school$school_type[1:3],df_school$name[1:3])
#> [1] "public"                  "public"
#> [3] "public"                  "Bethel Regional High School"
#> [5] "Ayagina'ar Elitnaurvik"  "Kwigillingok School"
```

3.3 Missing values

Missing values

Missing values have the value `NA`

`NA` is a special keyword, not the same as the character string `"NA"`

use `is.na()` function to determine if a value is missing

`is.na()` returns a logical vector

```r
is.na(5)
#> [1] FALSE
is.na(NA)
#> [1] TRUE
is.na("NA")
#> [1] FALSE
typeof(is.na("NA")) # example of a logical vector
#> [1] "logical"

nvector <- c(10,5,NA)
is.na(nvector)
#> [1] FALSE FALSE  TRUE
typeof(is.na(nvector)) # example of a logical vector
#> [1] "logical"

svector <- c("e","f",NA,"NA")
is.na(svector)
#> [1] FALSE FALSE  TRUE FALSE
```

Missing values are "contagious"

What does "contagious" mean?

operations involving a missing value will yield a missing value

```
7>5
#> [1] TRUE
7>NA
#> [1] NA
0==NA
#> [1] NA
2*c(0,1,2,NA)
#> [1]  0  2  4 NA
NA*c(0,1,2,NA)
#> [1] NA NA NA NA
```

Functions and missing values example, `table()`

`table()` function is useful for investigating categorical variables

```
str(df_event$event_type)
#>  chr [1:18680] "public hs" "public hs" "public hs" "public hs" ...
table(df_event$event_type)
#>
#> 2yr college 4yr college        other   private hs   public hs
#>         951          531         2001         3774       11423
```

Functions and missing values example, `table()`

By default `table()` ignores `NA` values

```
#?table
str(df_event$school_type_pri)
#>  int [1:18680] NA NA NA NA NA 1 1 NA 1 NA ...
table(df_event$school_type_pri)
#>
#>    1    2    5
#> 3765    8    1
```

`useNA` argument controls if table includes counts of `NA`s. Allowed values:

  never ("no") [DEFAULT VALUE]

  only if count is positive ("ifany");

  even for zero counts ("always")"

```
nrow(df_event)
#> [1] 18680
table(df_event$school_type_pri, useNA="always")
#>
#>    1    2    5  <NA>
#> 3765    8    1 14906
```

Broader point: Most functions that create descriptive statistics have options about how to treat missing values'

  When investigating data, good practice to *always* show missing values

4 Find homework groups

What we'll do to choose homework groups

Meet new people (10 minutes of speed-dating!)

find someone in class you don't know and talk to them for two minutes about anything

e.g., where you from, what program, what are research interests, what you like doing outside of school, work

Enrolled students choose homework groups (10 minutes)

one side of room for students who want to work collaboratively on problem sets

one side of room for students who want to work mostly on their own (e.g,. due to full work/family schedule);

must be groups of 3

cannot have more than 2 people from same academic program (e.g., HEOC, HDP)

Auditors not part of "official" homework groups of 3, but they are welcome to join any homework group or form their own homework group

Recommendation: **Use Zoom for group meetings!**

https://ucla.zoom.us/

5 Investigating data patterns [tidyverse]

Introduction to the `dplyr` library

`dplyr`, a package within the `tidyverse` suite of packages, provide tools for manipulating data frames

Wickham describes functions within `dplyr` as a set of "verbs" that fall in the broader categories of **subsetting**, **sorting**, and **transforming**

Today	Upcoming weeks
Subsetting data	**Transforming data**
- `select()` variables	- `mutate()` creates new variables
- `filter()` observations	- `summarize()` calculates across rows
Sorting data	- `group_by()` to calculate across rows within groups
- `arrange()`	

All `dplyr` verbs (i.e., functions) work as follows

1. first argument is a data frame
2. subsequent arguments describe what to do with variables and observations in data frame

    refer to variable names without quotes
3. result of the function is a new data frame

5.1 Select variables

Select variables using `select()` function

Printing observations is key to investigating data, but datasets often have hundreds, thousands of variables

`select()` function selects **columns** of data (i.e., variables) you specify

first argument is the name of data frame object

remaining arguments are variable names, which are separated by commas and without quotes

Without **assignment** (`<-`), `select()` by itself simply prints selected vars

```
#?select
select(df_event,instnm,event_date,event_type,event_state,med_inc)
#> # A tibble: 18,680 x 5
#>    instnm      event_date event_type event_state med_inc
#>    <chr>       <date>     <chr>      <chr>         <dbl>
#>  1 UM Amherst  2017-10-12 public hs  MA            71714.
#>  2 UM Amherst  2017-10-04 public hs  MA            89122.
#>  3 UM Amherst  2017-10-25 public hs  MA            70136.
#>  4 UM Amherst  2017-10-26 public hs  MA            70136.
#>  5 Stony Brook 2017-10-02 public hs  MA            71024.
#>  6 USCC        2017-09-18 private hs MA            71024.
#>  7 UM Amherst  2017-09-18 private hs MA            71024.
#>  8 UM Amherst  2017-09-26 public hs  MA            97225
#>  9 UM Amherst  2017-09-26 private hs MA            97225
#> 10 UM Amherst  2017-10-12 public hs  MA            77800.
#> # ... with 18,670 more rows
```

Select variables using `select()` function

Recall that all `dplyr` functions (e.g., `select()`) return a new data frame object

type equals "list"

length equals number of vars you select

```
typeof(select(df_event,instnm,event_date,event_type,event_state,med_inc))
#> [1] "list"
length(select(df_event,instnm,event_date,event_type,event_state,med_inc))
#> [1] 5
```

`glimpse()` : tidyverse function for viewing data frames

a cross between `str()` and simply printing data

```
?glimpse
glimpse(df_event)
```

`glimpse()` a `select()` set of variables

```
glimpse(select(df_event,instnm,event_date,event_type,event_state,med_inc))
#> Observations: 18,680
#> Variables: 5
#> $ instnm      <chr> "UM Amherst", "UM Amherst", "UM Amherst", "UM Amhe...
#> $ event_date  <date> 2017-10-12, 2017-10-04, 2017-10-25, 2017-10-26, 2...
#> $ event_type  <chr> "public hs", "public hs", "public hs", "public hs"...
#> $ event_state <chr> "MA", "MA", "MA", "MA", "MA", "MA", "MA", "MA", "M...
#> $ med_inc     <dbl> 71713.5, 89121.5, 70136.5, 70136.5, 71023.5, 71023...
```

Select variables using `select()` function

With **assignment** (`<-`), `select()` creates a new object containing only the variables you specify

```
event_small <- select(df_event,instnm,event_date,event_type,event_state,med_inc)

glimpse(event_small)
#> Observations: 18,680
#> Variables: 5
#> $ instnm      <chr> "UM Amherst", "UM Amherst", "UM Amherst", "UM Amhe...
#> $ event_date  <date> 2017-10-12, 2017-10-04, 2017-10-25, 2017-10-26, 2...
#> $ event_type  <chr> "public hs", "public hs", "public hs", "public hs"...
#> $ event_state <chr> "MA", "MA", "MA", "MA", "MA", "MA", "MA", "MA", "M...
#> $ med_inc     <dbl> 71713.5, 89121.5, 70136.5, 70136.5, 71023.5, 71023...
```

Select

`select()` can use "helper functions" `starts_with()`, `contains()`, and `ends_with()` to choose columns

```
?select()
```

Example:

```r
#names(df_event)

select(df_event,instnm,starts_with("event"))
#> # A tibble: 18,680 x 8
#>    instnm event_date event_type event_state event_inst event_name
#>    <chr>  <date>     <chr>      <chr>       <chr>      <chr>
#>  1 UM Am~ 2017-10-12 public hs  MA          In-State   Amherst-P~
#>  2 UM Am~ 2017-10-04 public hs  MA          In-State   Hampshire~
#>  3 UM Am~ 2017-10-25 public hs  MA          In-State   Chicopee ~
#>  4 UM Am~ 2017-10-26 public hs  MA          In-State   Chicopee ~
#>  5 Stony~ 2017-10-02 public hs  MA          Out-State  Easthampt~
#>  6 USCC   2017-09-18 private hs MA          Out-State  Williston~
#>  7 UM Am~ 2017-09-18 private hs MA          In-State   Williston~
#>  8 UM Am~ 2017-09-26 public hs  MA          In-State   Granby Jr~
#>  9 UM Am~ 2017-09-26 private hs MA          In-State   MacDuffie~
#> 10 UM Am~ 2017-10-12 public hs  MA          In-State   Smith Aca~
#> # ... with 18,670 more rows, and 2 more variables:
#> #   event_location_name <chr>, event_datetime_start <dttm>
```

Rename variables

rename() function renames variables within a data frame object

Syntax:

```
rename(obj_name, new_name = old_name,...)
```

```
rename(df_event, g12_offered = g12offered,
       titlei = titlei_status_pub)
names(df_event)
```

Variable names do not change permanently unless we combine rename with assignment

```
rename_event <- rename(df_event, g12_offered = g12offered, titlei = titlei_statu
names(rename_event)
rm(rename_event)
```

5.2 Filter rows

The `filter()` function

`filter()` allows you to **select observations** based on values of variables

Arguments

first argument is name of data frame

subsequent arguments are *logical expressions* to filter the data frame

Multiple expressions separated by commas work as **AND** operators (e.g., condtion 1 `TRUE` AND condition 2 `TRUE`)

What is the result of a `filter()` command?

`filter()` returns a data frame consisting of rows where the condition is `TRUE`

```
?filter
```

Example from data frame object `df_school`, each obs is a high school

Show all obs where the high school received 1 visit from UC Berkeley (110635) [output omitted]

```
filter(df_school,visits_by_110635 == 1)
```

Note that resulting object is list, consisting of obs where condition `TRUE`

```
nrow(df_school)
#> [1] 21301
nrow(filter(df_school,visits_by_110635 == 1))
#> [1] 528
```

Filter, character variables

Use single quotes `''` or double quotes `""` to refer to values of character variables

```
glimpse(select(df_school, school_type,state_code))
#> Observations: 21,301
#> Variables: 2
#> $ school_type <chr> "public", "public", "public", "public", "public", ...
#> $ state_code  <chr> "AK", "AK", "AK", "AK", "AK", "AK", "AK", "AK", "A...
```

Identify all private high schools in CA that got 1 visit by particular universities

UC-Berkeley (ID=110635)

```
filter(df_school,visits_by_110635 == 1, school_type == "private",
       state_code == "CA")
```

University of Alabama (ID=100751)

```
filter(df_school,visits_by_100751 == 1, school_type == "private",
       state_code == "CA")
```

Visited once by Berkeley and University of Alabama

```
filter(df_school,visits_by_100751 == 1, visits_by_110635 == 1,
       school_type == "private", state_code == "CA")
```

Logical operators for comparisons

Symbol	Meaning
==	Equal to
!=	Not equal to
>	greater than
>=	greater than or equal to
<	less than
<=	less than or equal to
&	AND
|	OR
%in	includes

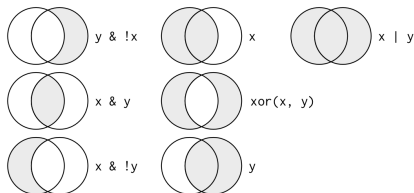
Figure 1: "Boolean" operations, x=left circle, y=right circle, from Wichkam (2018)

Filters and comparisons, Demonstration

Schools visited by Bama (`100751`) and/or Berkeley (`110635`)

```
#berkeley and bama
filter(df_school,visits_by_100751 >= 1, visits_by_110635 >= 1)
filter(df_school,visits_by_100751 >= 1 & visits_by_110635 >= 1) # same same
#berkeley or bama
filter(df_school,visits_by_100751 >= 1 | visits_by_110635 >= 1)
```

Apply `count()` function on top of `filter()` function to count the number of observations that satisfy criteria

Avoids printing individual observations

```
#Number of schools that git visit by Berkeley AND Bama
count(filter(df_school,visits_by_100751 >= 1 & visits_by_110635 >= 1))
#> # A tibble: 1 x 1
#>        n
#>    <int>
#> 1    247
#Number of schools that git visit by Berkeley OR Bama
count(filter(df_school,visits_by_100751 >= 1 | visits_by_110635 >= 1))
#> # A tibble: 1 x 1
#>        n
#>    <int>
#> 1   2763
```

Filters and comparisons, >=

Number of public high schools that are at least 50% Black in Alabama compared to number of schools that received visit by Bama

```
#at least 50% black
count(filter(df_school, school_type == "public", pct_black >= 50,
             state_code == "AL"))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1    86
count(filter(df_school, school_type == "public", pct_black >= 50,
             state_code == "AL", visits_by_100751 >= 1))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1    21
#at least 50% white
count(filter(df_school, school_type == "public", pct_white >= 50,
             state_code == "AL"))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1   238
count(filter(df_school, school_type == "public", pct_white >= 50,
             state_code == "AL", visits_by_100751 >= 1))
#> # A tibble: 1 x 1
#>       n
```

Filters and comparisons, not equals (!=)

Count the number of high schools visited by University of Colorado (126614) that are not located in CO

```
#number of high schools visited by U Colorado
count(filter(df_school, visits_by_126614 >= 1))
#> # A tibble: 1 x 1
#>         n
#>     <int>
#> 1  1056

#number of high schools visited by U Colorado not located in CO
count(filter(df_school, visits_by_126614 >= 1, state_code != "CO"))
#> # A tibble: 1 x 1
#>         n
#>     <int>
#> 1    873
#number of high schools visited by U Colorado located in CO
#count(filter(df_school, visits_by_126614 >= 1, state_code == "CO"))
```

Filters and comparisons, %in% operator

What if you wanted to count the number of schools visited by Bama (100751) in a group of states?

```
count(filter(df_school,visits_by_100751 >= 1, state_code == "MA" |
              state_code == "VT" | state_code == "ME"))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1   108
```

Easier way to do this is with `%in%` operator

```
count(filter(df_school,visits_by_100751 >= 1, state_code %in% c("MA","ME","VT"))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1   108
```

Select the private high schools that got either 2 or 3 visits from Bama

```
count(filter(df_school, visits_by_100751 %in% 2:3, school_type == "private"))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1   183
```

Identifying data type and possible values helpful for filtering

`class()` and `str()` shows data type of a variable

`table()` to show potential values of categorical variables

```
class(df_event$event_type)
#> [1] "character"
str(df_event$event_type)
#>  chr [1:18680] "public hs" "public hs" "public hs" "public hs" ...
table(df_event$event_type, useNA="always")
#>
#> 2yr college 4yr college        other  private hs    public hs        <NA>
#>         951         531         2001        3774       11423           0
class(df_event$event_state)
#> [1] "character"
str(df_event$event_state) # double quotes indicate character
#>  chr [1:18680] "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" "MA" ...
class(df_event$med_inc)
#> [1] "numeric"
str(df_event$med_inc)
#>  num [1:18680] 71714 89122 70137 70137 71024 ...
```

Now that we know `event_type` is a character, we can filter values

```
count(filter(df_event, event_type == "public hs", event_state =="CA"))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1  1100
```

Filtering and missing values

Wickham (2018) states:

> " `filter()` only includes rows where condition is TRUE; it excludes both `FALSE` and `NA` values. To preserve missing values, ask for them explicitly."

Investigate var `df_event$fr_lunch` , number of free/reduced lunch students

only available for visits to public high schools

```
#visits to public HS with less than 50 students on free/reduced lunch
count(filter(df_event,event_type == "public hs", fr_lunch<50))
#> # A tibble: 1 x 1
#>        n
#>    <int>
#> 1   910
#visits to public HS, where free/reduced lunch missing
count(filter(df_event,event_type == "public hs", is.na(fr_lunch)))
#> # A tibble: 1 x 1
#>         n
#>     <int>
#> 1    26
#visits to public HS, where free/reduced is less than 50 OR is missing
count(filter(df_event,event_type == "public hs", fr_lunch<50 | is.na(fr_lunch)))
#> # A tibble: 1 x 1
#>         n
#>     <int>
#> 1    936
```

Exercise

Task

Create a filter to identify all the high schools that recieved 1 visit from UC Berkeley (110635) AND 1 visit from CU Boulder (126614)[output omitted]

Solution

```
filter(df_school,visits_by_110635 == 1, visits_by_126614==1)

nrow(filter(df_school,visits_by_110635 == 1, visits_by_126614==1))
count(filter(df_school,visits_by_110635 == 1, visits_by_126614==1))
```

Must **assign** to create new object based on filter

```
berk_boulder <- filter(df_school,visits_by_110635 == 1, visits_by_126614==1)
count(berk_boulder)
```

Exercises

Use the data from df_event, which has one observation for each off-campus recruiting event a university attends

1. Count the number of events attended by the University of Pittsburgh (Pitt) `univ_id == 215293`

2. Count the number of recruiting events by Pitt at public or private high schools

3. Count the number of recruiting events by Pitt at public or private high schools located in the state of PA

4. Count the number of recruiting events by Pitt at public high schools not located in PA where median income is less than 100,000

5. Count the number of recruiting events by Pitt at public high schools not located in PA where median income is greater than or equal to 100,000

6. Count the number of out-of-state recruiting events by Pitt at private high schools or public high schools with median income of at least 100,000

Solution

1. Count the number of events attended by the University of Pittsburgh (Pitt)
   `univ_id == 215293`

```
count(filter(df_event, univ_id == 215293))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1  1225
```

2. Count the number of recruiting events by Pitt at public or private high
   schools

```
str(df_event$event_type)
#>  chr [1:18680] "public hs" "public hs" "public hs" "public hs" ...
table(df_event$event_type, useNA = "always")
#>
#> 2yr college  4yr college       other  private hs   public hs        <NA>
#>         951          531        2001        3774       11423           0
count(filter(df_event, univ_id == 215293, event_type == "private hs" |
                 event_type == "public hs"))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1  1030
```

Solution

3. Count the number of recruiting events by Pitt at public or private high schools located in the state of PA

```
count(filter(df_event, univ_id == 215293, event_type == "private hs" |
                event_type == "public hs", event_state == "PA"))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1   262
```

4. Count the number of recruiting events by Pitt at public high schools not located in PA where median income is less than 100,000

```
count(filter(df_event, univ_id == 215293, event_type == "public hs",
              event_state != "PA", med_inc < 100000))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1   213
```

Solution

5. Count the number of recruiting events by Pitt at public high schools not located in PA where median income is greater than or equal to 100,000

```
count(filter(df_event, univ_id == 215293, event_type == "public hs",
             event_state != "PA", med_inc >= 100000))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1   344
```

6. Count the number of out-of-state recruiting events by Pitt at private high schools or public high schools with median income of at least 100,000

```
count(filter(df_event, univ_id == 215293, event_state != "PA",
             (event_type == "public hs" & med_inc >= 100000) |
              event_type == "private hs"))
#> # A tibble: 1 x 1
#>       n
#>   <int>
#> 1   553
```

5.3 Arrange rows

arrange() function

arrange() function "arranges" rows in a data frame; said different, it sorts observations

Syntax: `arrange(x,...)`

> First argument, `x`, is a data frame
>
> Subsequent arguments are a "comma separated list of unquoted variable names"

```
arrange(df_event, event_date)
```

Data frame goes back to previous order unless you **assign** the new order

```
df_event
df_event <- arrange(df_event, event_date)
df_event
```

arrange() function

Ascending and descending order

arrange() sorts in **ascending** order by default

use desc() to sort a column by descending order

arrange(df_event, **desc**(event_date))

Can sort by multiple variables

arrange(df_event, univ_id, **desc**(event_date), **desc**(med_inc))

#sort by university and descending by size of 12th grade class; combine with sel
select(**arrange**(df_event, univ_id, **desc**(g12)),instnm,event_type,event_date,g12)

`arrange()`, missing values sorted at the end

Missing values automatically sorted at the end, regardless of whether you sort ascending or descending

Below, we sort by university, then by date of event, then by ID of high school

```
#by university, date, ascending school id
select(arrange(df_event, univ_id, desc(event_date), school_id),
       instnm,event_date,event_type,school_id)

#by university, date, descending school id
select(arrange(df_event, univ_id, desc(event_date), desc(school_id)),
       instnm,event_date,event_type,school_id)
```

Can sort by `is.na` to put missing values first

```
select(arrange(df_event, univ_id, desc(event_date), desc(is.na(school_id))),
       instnm,event_date,event_type,school_id)
#> # A tibble: 18,680 x 4
#>    instnm event_date event_type school_id
#>    <chr>  <date>     <chr>      <chr>
#>  1 Bama   2017-12-18 other      <NA>
#>  2 Bama   2017-12-18 private hs A9106483
#>  3 Bama   2017-12-15 other      <NA>
#>  4 Bama   2017-12-15 public hs  484473005095
#>  5 Bama   2017-12-15 public hs  062927004516
#>  6 Bama   2017-12-14 other      <NA>
#>  7 Bama   2017-12-13 other      <NA>
#>  8 Bama   2017-12-13 public hs  130387001439
#>  9 Bama   2017-12-13 private hs 00071151
```

Exercise, arranging

Use the data from df_event, which has one observation for each off-campus recruiting event a university attends

1. Sort ascending by "univ_id" and descending by "event_date"
2. Select four variables in total and sort ascending by "univ_id" and descending by "event_date"
3. Now using the same variables from above, sort by `is.na` to put missing values in "school_id" first

Solution

1. Sort ascending by "univ_id" and descending by "event_date"

```
arrange(df_event, univ_id, desc(event_date))
#> # A tibble: 18,680 x 33
#>    instnm univ_id instst   pid event_date event_type  zip    school_id
#>    <chr>    <int> <chr> <int> <date>     <chr>       <chr>  <chr>
#>  1 Bama    100751 AL     7115 2017-12-18 private hs  77089  A9106483
#>  2 Bama    100751 AL     7121 2017-12-18 other       <NA>   <NA>
#>  3 Bama    100751 AL     7114 2017-12-15 public hs   75165  48447300~
#>  4 Bama    100751 AL     7100 2017-12-15 public hs   93012  06292700~
#>  5 Bama    100751 AL     7073 2017-12-15 other       98027  <NA>
#>  6 Bama    100751 AL     7072 2017-12-14 other       98007  <NA>
#>  7 Bama    100751 AL     7118 2017-12-13 public hs   31906  13038700~
#>  8 Bama    100751 AL     7099 2017-12-13 private hs  90293  00071151
#>  9 Bama    100751 AL     7109 2017-12-13 public hs   92630  06338600~
#> 10 Bama    100751 AL     7071 2017-12-13 other       98032  <NA>
#> # ... with 18,670 more rows, and 25 more variables: ipeds_id <int>,
#> #   event_state <chr>, event_inst <chr>, med_inc <dbl>, pop_total <dbl>,
#> #   pct_white_zip <dbl>, pct_black_zip <dbl>, pct_asian_zip <dbl>,
#> #   pct_hispanic_zip <dbl>, pct_amerindian_zip <dbl>,
#> #   pct_nativehawaii_zip <dbl>, pct_tworaces_zip <dbl>,
#> #   pct_otherrace_zip <dbl>, fr_lunch <dbl>, titlei_status_pub <fct>,
#> #   total_12 <dbl>, school_type_pri <int>, school_type_pub <int>,
#> #   g12offered <dbl>, g12 <dbl>, total_students_pub <dbl>,
#> #   total_students_pri <dbl>, event_name <chr>, event_location_name <chr>,
#> #   event_datetime_start <dttm>
```

Solution

2. Select four variables in total and sort ascending by "univ_id" and descending by "event_date"

```
select(arrange(df_event, univ_id, desc(event_date)), univ_id, event_date,
       instnm, event_type)
#> # A tibble: 18,680 x 4
#>    univ_id event_date instnm event_type
#>      <int> <date>     <chr>  <chr>
#>  1  100751 2017-12-18 Bama   private hs
#>  2  100751 2017-12-18 Bama   other
#>  3  100751 2017-12-15 Bama   public hs
#>  4  100751 2017-12-15 Bama   public hs
#>  5  100751 2017-12-15 Bama   other
#>  6  100751 2017-12-14 Bama   other
#>  7  100751 2017-12-13 Bama   public hs
#>  8  100751 2017-12-13 Bama   private hs
#>  9  100751 2017-12-13 Bama   public hs
#> 10  100751 2017-12-13 Bama   other
#> # ... with 18,670 more rows
```

Solution

3. Select the variables "univ_id", "event_date", and "school_id" and sort by `is.na` to put missing values in "school_id" first.

```
select(arrange(df_event, univ_id, desc(event_date), desc(is.na(school_id))),
       univ_id, event_date, school_id)
#> # A tibble: 18,680 x 3
#>    univ_id event_date school_id
#>      <int> <date>     <chr>
#>  1  100751 2017-12-18 <NA>
#>  2  100751 2017-12-18 A9106483
#>  3  100751 2017-12-15 <NA>
#>  4  100751 2017-12-15 484473005095
#>  5  100751 2017-12-15 062927004516
#>  6  100751 2017-12-14 <NA>
#>  7  100751 2017-12-13 <NA>
#>  8  100751 2017-12-13 130387001439
#>  9  100751 2017-12-13 00071151
#> 10  100751 2017-12-13 063386005296
#> # ... with 18,670 more rows
```

6 Appendix: directories and filepaths [for your reference]

Working directory

(Current) Working directory

the folder/directory in which you are currently working

this is where R looks for files

Files located in your current working directory can be accessed without specifying a filepath because R automatically looks in this folder

Function `getwd()` shows current working directory

```
getwd()
#> [1] "C:/Users/ozanj/Documents/rclass/lectures/lecture2"
```

Command `list.files()` lists all files located in working directory

```
getwd()
#> [1] "C:/Users/ozanj/Documents/rclass/lectures/lecture2"
list.files()
#>  [1] "fp1.JPG"                      "fp2.JPG"
#>  [3] "lecture2.1_ucla.pdf"          "lecture2.1_ucla.Rmd"
#>  [5] "lecture2.1_ucla.tex"          "lecture2.pdf"
#>  [7] "lecture2.Rmd"                 "lecture2.tex"
#>  [9] "lecture2_test.Rmd"            "lecture2_test.tex"
#> [11] "lecture2_ucla.pdf"            "lecture2_ucla.Rmd"
#> [13] "lecture2_ucla.tex"            "problemset2.html"
#> [15] "problemset2.Rmd"              "problemset2_solutions.html"
#> [17] "problemset2_solutions.html.zip" "problemset2_solutions.Rmd"
#> [19] "sample.html"                  "sample.Rmd"
#> [21] "sample_simple_rmarkdown.txt"  "text"
```

Working directory, "Code chunks" vs. "console" and "R scripts"

When you run **code chunks** in RMarkdown files (.Rmd), the working directory is set to the filepath where the .Rmd file is stored

```
getwd()
#> [1] "C:/Users/ozanj/Documents/rclass/lectures/lecture2"
list.files()
#>  [1] "fp1.JPG"                        "fp2.JPG"
#>  [3] "lecture2.1_ucla.pdf"            "lecture2.1_ucla.Rmd"
#>  [5] "lecture2.1_ucla.tex"            "lecture2.pdf"
#>  [7] "lecture2.Rmd"                   "lecture2.tex"
#>  [9] "lecture2_test.Rmd"              "lecture2_test.tex"
#> [11] "lecture2_ucla.pdf"              "lecture2_ucla.Rmd"
#> [13] "lecture2_ucla.tex"              "problemset2.html"
#> [15] "problemset2.Rmd"                "problemset2_solutions.html"
#> [17] "problemset2_solutions.html.zip" "problemset2_solutions.Rmd"
#> [19] "sample.html"                    "sample.Rmd"
#> [21] "sample_simple_rmarkdown.txt"    "text"
#> [23] "transform-logical.png"
```

When you run code from the **R Console** or an **R Script**, the working directory is....

Command `getwd()` shows current working directory

```
getwd()
#> [1] "C:/Users/ozanj/Documents/rclass/lectures/lecture2"
```

Absolute vs. relative filepath

Absolute file path: The absolute file path is the complete list of directories needed to locate a file or folder.

```
setwd("/Users/pm/Desktop/rclass/lectures/lecture2")
```

Relative file path: The relative file path is the path relative to your current location/directory. Assuming your current working directory is in the "lecture2" folder and you want to change your directory to the data folder, your relative file path would look something like this:

```
setwd("../../data")
```

File path shortcuts

| Key	Description
~	tilde is a shortcut for user's home directory (mine is my name pm)
../	moves up a level
../../	moves up two level

Exercise

1. Let's create a folder on our desktop and name it red
2. Inside the red folder, create two subfolders named orange and yellow
3. Inside the yellow folder create another subfolder named green

Make sure to name these folders in lowercase.

You should have 1 folder on your desktop called red. Inside the red folder you have two folders called orange and yellow. Inside the yellow folder you have a folder called green.

Here is a visual of how it should look...

File path visual
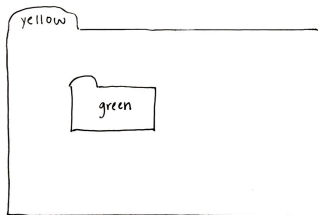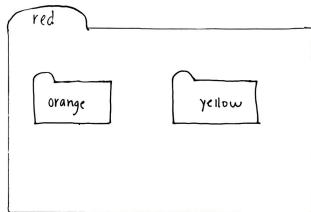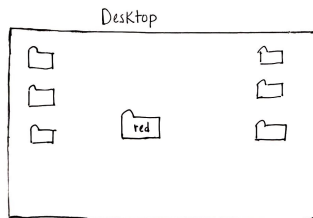
Exercise continued

Let's say we want to get to the green folder using the absolute file path.

1. View your current working directory `getwd()`

2. Set your working directory to the green folder using the absolute file path

3. Now set your working directory to the orange folder using the relative file path (hint: ../)

Solution

```r
getwd()
setwd("~/Desktop/red/yellow/green")
getwd()
setwd("../../orange")
getwd()
```