

## Lecture 5: Processing across rows

### Managing and Manipulating Data Using R

# Introduction

## Required reading for next week:

- ▶ Grolemund and Wickham 5.6 - 5.7 (grouped summaries and mutates)
- ▶ Xie, Allaire, and Grolemund 4.1 (R Markdown, ioslides presentations) [LINK HERE](#) and 4.3 (R Markdown, Beamer presentations) [LINK HERE](#)
  - ▶ Why? Lectures for this class are `beamer_presentation` output type.
  - ▶ `ioslides_presentation` are the most basic presentation output format for RMarkdown, so learning about `ioslides` will help you understand `beamer`

# What we will do today

1. Introduction
2. Introduce `group_by()` and `summarize()`
  - 2.1 `group_by`
  - 2.2 `summarize()`
3. Combining `group_by()` and `summarize()`
  - 3.1 `summarize()` and Counts
  - 3.2 `summarize()` and means
  - 3.3 `summarize()` and logical vectors, part II
4. Summarize multiple columns
5. Attach aggregate measures to your data frame

## Libraries we will use today

“Load” the package we will use today (output omitted)

▶ **you must run this code chunk**

```
library(tidyverse)
```

If package not yet installed, then must install before you load. Install in “console” rather than .Rmd file

▶ Generic syntax: `install.packages("package_name")`

▶ Install “tidyverse”: `install.packages("tidyverse")`

Note: when we load package, name of package is not in quotes; but when we install package, name of package is in quotes:

▶ `install.packages("tidyverse")`

▶ `library(tidyverse)`

# Data we will use today

Data on off-campus recruiting events by public universities

► Object `df_event`

► One observation per university, recruiting event

```
rm(list = ls()) # remove all objects
```

```
#load dataset with one obs per recruiting event
```

```
load(url("https://github.com/ozanj/rclass/raw/master/data/recruiting/recruit_ev
```

```
#load("../..../data/recruiting/recruit_event_allvars.Rdata")
```

## Processing across variables vs. processing across observations

Visits by UC Berkeley to public high schools

```
#> # A tibble: 5 x 6
#>   school_id    state tot_stu_pub fr_lunch pct_fr_lunch med_inc
#>   <chr>      <chr>      <dbl>    <dbl>      <dbl>    <dbl>
#> 1 340882002126 NJ          1846      29      0.0157 178732
#> 2 340147000250 NJ          1044      50      0.0479  62288
#> 3 340561003796 NJ          1505     298      0.198 100684.
#> 4 340165005124 NJ          1900      43      0.0226 160476.
#> 5 341341003182 NJ          1519     130      0.0856 144346
```

So far, we have focused on “processing across variables”

- ▶ Performing calculations across columns (i.e., vars), typically within a row (i.e., observation)
- ▶ Example: percent free-reduced lunch (above)

Processing across obs (focus of today's lecture)

- ▶ Performing calculations across rows (i.e., obs), often within a column (i.e., variable)
- ▶ Example: Average household income of visited high schools, by state

# Why processing across observations

## Note

- ▶ in today's lecture, I'll use the terms "observations" and "rows" interchangeably

## Creation of analysis datasets often requires calculations across obs

Examples:

- ▶ You have a dataset with one observation per student-term and want to create a variable of credits attempted per term
- ▶ You have a dataset with one observation per student-term and want to create a variable of GPA for the semester or cumulative GPA for all semesters
- ▶ Number of off-campus recruiting events university makes to each state
- ▶ Average household income at visited versus non-visited high schools

## Creating graphs and tables of descriptive stats usually require calculations across obs

Example: Want to create a graph that shows number of recruiting events by event "type" (e.g., public HS, private HS) for each university

- ▶ Start with `df_event` dataset that has one observation per university, recruiting event
- ▶ Create new data frame object that has one observation per university and event type and has variable for number of events
  - ▶ this variable calculated by counting number of rows in each combination of university and event type
- ▶ This new data frame object is the input for creating desired graph



Introduce `group_by()` and `summarize()`

# Strategy for teaching processing across obs

In `tidyverse` the `group_by()` and `summarize()` functions are the primary means of performing calculations across observations

- ▶ Usually, processing across observations requires using `group_by()` and `summarize()` together
- ▶ `group_by()` typically not very useful by itself
- ▶ `summarize()` [with or without `group_by()` ] can be helpful for creating summary statistics that are the inputs for tables or graphs you create

How we'll teach:

- ▶ introduce `group_by()` and `summarize()` separately
  - ▶ goal: you understand what each function does
- ▶ then we'll combine them

group\_by

## group\_by()

**Description:** “ `group_by()` takes an existing data frame and converts it into a grouped data frame where operations are performed “by group”. `ungroup()` removes grouping.”

- ▶ part of **dplyr** package within **tidyverse**; not part of **Base R**
- ▶ works best with pipes `%>%` and `summarize()` function [described below]

**Basic syntax:** `group_by(.data, ...)`

- ▶ `.data` argument refers to name of data frame
- ▶ `...` argument refers to names of “group\_by” variables, separated by commas
  - ▶ Can “group by” one or many variables
  - ▶ Typically, “group\_by” variables are character, factor, or integer variables

Possible “group by” variables in `df_event` data:

- ▶ university name/id; event type (e.g., public HS, private HS); state

**Example:** in `df_event` , create frequency count of `event_type` [output omitted]

```
names(df_event)
#without group_by()
df_event %>% count(event_type)
df_event %>% count(instnm)
#group_by() university
df_event %>% group_by(instnm) %>% count(event_type)
```

## group\_by()

By itself `group_by()` doesn't do much; it just prints data

► Below, group `df_event` data by university, event type, and event state

```
#print object
df_event
#group_by (without pipes)
group_by(df_event, univ_id, event_type, event_state)
#group_by (with pipes)
df_event %>% group_by(univ_id, event_type, event_state)
df_event %>% group_by(univ_id, event_type, event_state) %>% glimpse()
```

But once an object is grouped, all subsequent functions are run separately “by group”

► recall that `count()` counts number of observations by group

```
# count number of observations in group, ungrouped data
df_event %>% count()
#group by and then count obs
df_event %>% group_by(univ_id) %>% count()
df_event %>% group_by(univ_id) %>% count() %>% glimpse()

df_event %>% group_by(univ_id, event_type) %>% count()
df_event %>% group_by(univ_id, event_type) %>% count() %>% glimpse()

df_event %>% group_by(univ_id, event_type, event_state) %>% count()
df_event %>% group_by(univ_id, event_type, event_state) %>% count() %>% glimpse()
```

## Grouping not retained unless you **assign** it

Below, we'll use `class()` function to show whether data frame is grouped

- ▶ will talk more about `class()` next week, but for now, just think of it as a function that provides information about an object
- ▶ similar to `typeof()`, but `class()` provides different info about object

Grouping is not retained unless you **assign** it

```
class(df_event)
#> [1] "tbl_df"      "tbl"        "data.frame"
df_event_grp <- df_event %>% group_by(univ_id, event_type, event_state) # using
class(df_event_grp)
#> [1] "grouped_df" "tbl_df"     "tbl"        "data.frame"
```

Use `ungroup(object)` to un-group grouped data

```
class(df_event_grp)
#> [1] "grouped_df" "tbl_df"     "tbl"        "data.frame"
df_event_grp <- ungroup(df_event_grp)
class(df_event_grp)
#> [1] "tbl_df"      "tbl"        "data.frame"
rm(df_event_grp)
```

## `group_by()` student exercise

1. Group by “instnm” and get a frequency count.
  - ▶ How many rows and columns do you have? What do the number of rows mean?
2. Now group by “instnm” **and** “event\_type” and get a frequency count.
  - ▶ How many rows and columns do you have? What do the number of rows mean?
3. **Bonus:** In the same code chunk, group by “instnm” and “event\_type”, but this time filter for observations where “med\_inc” is greater than 75000 and get a frequency count.

## group\_by() student exercise solutions

1. Group by "instnm" and get a frequency count.

► How many rows and columns do you have? What do the number of rows mean?

```
df_event %>%  
  group_by(instnm) %>%  
  count()  
#> # A tibble: 16 x 2  
#>   instnm      n  
#>   <chr>    <int>  
#> 1 Arkansas    994  
#> 2 Bama      4258  
#> 3 Cinci      679  
#> 4 CU Boulder 1439  
#> 5 Kansas    1014  
#> 6 NC State   640  
#> 7 Pitt      1225  
#> 8 Rutgers   1135  
#> 9 S Illinois  549  
#> 10 Stony Brook 730  
#> 11 UC Berkeley 879  
#> 12 UC Irvine   539  
#> 13 UGA        827  
#> 14 UM Amherst  908  
#> 15 UNL       1397  
#> 16 USCC      1467
```



2. Now group by “instnm” **and** “event\_type” and get a frequency count.  
▶ How many rows and columns do you have? What do the number of rows mean?

```
df_event %>%  
  group_by(instnm, event_type) %>%  
  count()
```

```
#> # A tibble: 80 x 3  
#>   instnm event_type      n  
#>   <chr>   <chr>      <int>  
#> 1 Arkansas 2yr college    32  
#> 2 Arkansas 4yr college    14  
#> 3 Arkansas other        112  
#> 4 Arkansas private hs    222  
#> 5 Arkansas public hs    614  
#> 6 Bama     2yr college    127  
#> 7 Bama     4yr college    158  
#> 8 Bama     other        608  
#> 9 Bama     private hs    963  
#> 10 Bama    public hs    2402  
#> # ... with 70 more rows
```

## group\_by() student exercise solutions

3. **Bonus:** Group by “instnm” and “event\_type”, but this time filter for observations where “med\_inc” is greater than 75000 and get a frequency count.

```
df_event %>%  
  group_by(instnm, event_type) %>%  
  filter(med_inc > 75000) %>%  
  count()
```

```
#> # A tibble: 80 x 3  
#>   instnm event_type      n  
#>   <chr>   <chr>      <int>  
#> 1 Arkansas 2yr college      7  
#> 2 Arkansas 4yr college      3  
#> 3 Arkansas other          30  
#> 4 Arkansas private hs     99  
#> 5 Arkansas public hs    303  
#> 6 Bama     2yr college     21  
#> 7 Bama     4yr college     42  
#> 8 Bama     other         249  
#> 9 Bama     private hs    477  
#> 10 Bama    public hs    1478  
#> # ... with 70 more rows
```

summarize()

## summarize() function

**Description:** `summarize()` calculates across rows; then collapses into single row

- ▶ `summarize()` create scalar vars summarizing variables of existing data frame
- ▶ if you first group data frame using `group_by()`, `summarize()` creates summary vars separately for each group, returning object with one row per group
- ▶ if data frame not grouped, `summarize()` will result in one row.

**Syntax:** `summarize(.data, ...)`

- ▶ `.data`: a data frame; omit if using `summarize()` after pipe `%>%`
- ▶ `...`: Name-value pairs of summary functions separated by commas
  - ▶ “name” will be the name of new variable you will create
  - ▶ “value” should be expression that returns a single value like `min(x)`, `n()`
  - ▶ variable names do not need to be placed within quotes

**Value** (what `summarize()` returns/creates)

- ▶ Object of same class as `.data.`; object will have one obs per “by group”

**Useful functions (i.e., “helper functions”)**

- ▶ Standalone functions called *within* `summarize()`, e.g., `mean()`, `n()`
- ▶ e.g., count function `n()` takes no arguments; returns number of rows in group

**Example:** Count total number of events (output omitted)

```
summarize(df_event, num_events=n()) # without pipes
df_event %>% summarize(num_events=n()) # using pipes
```

## Investigate objects created by `summarize()`

**Example:** Count total number of events

```
df_event %>% summarize(num_events=n())  
df_event %>% summarize(num_events=n()) %>% str()
```

**Example:** What is max value of `med_inc` across all events

```
df_event %>% summarize(max_inc=max(med_inc, na.rm = TRUE))  
df_event %>% summarize(max_inc=max(med_inc, na.rm = TRUE)) %>% str()
```

**Example:** Count total number of events AND max value of median income

```
df_event %>% summarize(num_events=n(),  
                        max_inc=max(med_inc, na.rm = TRUE))  
df_event %>% summarize(num_events=n(),  
                        max_inc=max(med_inc, na.rm = TRUE)) %>% str()
```

### Takeaways

- ▶ by default, objects created by `summarize()` are data frames that contain variables created within `summarize()` and one observation [per “by group”]
- ▶ most “helper” functions (e.g., `max()`, `mean()`) have option `na.rm` to keep/remove missing obs before performing calculations
  - ▶ `na.rm = FALSE` (default); don't remove `NA`s prior to calculation
    - ▶ if any obs missing, then result of calculation is `NA`
  - ▶ `na.rm = TRUE` (default); remove `NA`s prior to calculation

## Retaining objects created by `summarize()`

Object created by `summarize()` not retained unless you **assign** it

```
event_temp <- df_event %>% summarize(num_events=n(),  
  mean_inc=mean(med_inc, na.rm = TRUE))
```

```
event_temp
```

```
#> # A tibble: 1 x 2
```

```
#>   num_events mean_inc
```

```
#>   <int>     <dbl>
```

```
#> 1     18680  89089.
```

```
rm(event_temp)
```

## Using `[]` operator to filter observations within summarize

Imagine we want to calculate avg. income, separately for in-state vs. out-of-state visits

► first, let's use `filter()` to make sure we can identify in-state vs. out-of-state

```
#in state
df_event %>% filter(event_state == instst) %>% count() %>% as_vector()
#>      n
#> 5425
#out state
df_event %>% filter(event_state != instst) %>% count() %>% as_vector()
#>      n
#> 13255
```

► calculate mean income for: all events; in-state events; out-of-state events

```
df_event %>%
  summarize(avg_inc = mean(med_inc, na.rm = TRUE), # all events
            avg_inc_inst = mean(med_inc[event_state == instst], na.rm = TRUE), # in-sta
            avg_inc_outst = mean(med_inc[event_state != instst], na.rm = TRUE) # out-st
  )
#> # A tibble: 1 x 3
#>   avg_inc avg_inc_inst avg_inc_outst
#>   <dbl>     <dbl>     <dbl>
#> 1  89089.      71589.      96162.
```

## Using `summarize()` to create descriptive statistics table

Often helpful to use `summarize()` to calculate summary statistics that are the basis for a table of descriptive statistics

**Task:** create a table of descriptive statistics about variable `med_inc`

► want these measures: number of non-missing obs; mean; standard deviation

```
df_event %>% mutate(non_miss_inc = is.na(med_inc)==0) %>%  
  summarize(  
    n = sum(non_miss_inc, na.rm = TRUE), #SAMPLE SIZE all  
    avg_inc = mean(med_inc, na.rm = TRUE), # MEAN  
    std_inc = sd(med_inc, na.rm = TRUE), # STANDARD DEVIATION all events  
  )
```

**Task:** same as above but separate measures for: all events; in-state; out-of-state

```
df_event %>% mutate(non_miss_inc = is.na(med_inc)==0) %>%  
  summarize(  
    n = sum(non_miss_inc, na.rm = TRUE), #SAMPLE SIZE  
    n_inst = sum(non_miss_inc[event_state == instst], na.rm = TRUE),  
    n_outst = sum(non_miss_inc[event_state != instst], na.rm = TRUE),  
    avg_inc = mean(med_inc, na.rm = TRUE), # MEAN  
    avg_inc_inst = mean(med_inc[event_state == instst], na.rm = TRUE),  
    avg_inc_outst = mean(med_inc[event_state != instst], na.rm = TRUE),  
    std_inc = sd(med_inc, na.rm = TRUE), # STANDARD DEVIATION  
    std_inc_inst = sd(med_inc[event_state == instst], na.rm = TRUE),  
    std_inc_outst = sd(med_inc[event_state != instst], na.rm = TRUE)  
  )
```



## `summarize()` student exercise

1. What is the min value of `med_inc` across all events?
  - ▶ Hint: Use `min()`
2. What is the mean value of `fr_lunch` across all events?
  - ▶ Hint: Use `mean()`

## summarize() student exercise

1. What is min value of `med_inc` across all events?

```
df_event %>%  
  summarize(min_med_income = min(med_inc, na.rm = TRUE))  
#> # A tibble: 1 x 1  
#>   min_med_income  
#>       <dbl>  
#> 1         12894.
```

## summarize() student exercise

2. What is the mean value of `fr_lunch` across all events?

► Hint: Use `mean()`

```
df_event %>%  
  summarize(mean_fr_lunch = mean(fr_lunch, na.rm = TRUE))  
#> # A tibble: 1 x 1  
#>   mean_fr_lunch  
#>       <dbl>  
#> 1         475.
```

Combining `group_by()` and `summarize()`

## Combining `summarize()` and `group_by`

`summarize()` on ungrouped vs. grouped data:

- ▶ By itself, `summarize()` performs calculations across all rows of data frame then collapses the data frame to a single row
- ▶ When data frame is grouped, `summarize()` performs calculations across rows within a group and then collapses to a single row for each group

**Example:** Count the number of events for each university

- ▶ remember: `df_event` has one observation per university, recruiting event

```
df_event %>% summarize(num_events=n())  
df_event %>% group_by(instnm) %>% summarize(num_events=n())
```

- ▶ Investigate the object created above

```
df_event %>% group_by(instnm) %>% summarize(num_events=n()) %>% str()
```

- ▶ Or we could retain object for later use

```
event_by_univ <- df_event %>% group_by(instnm) %>% summarize(num_events=n())  
str(event_by_univ)  
event_by_univ # print  
rm(event_by_univ)
```

## Combining `summarize()` and `group_by`

### Task

- ▶ Count number of recruiting events by institution and event\_type

```
df_event %>% group_by(instnm, event_type) %>% summarize(num_events=n())
```

*#investigate object created*

```
df_event %>% group_by(instnm, event_type) %>% summarize(num_events=n()) %>% str
```

Note that data frame object created by `group_by()` and `summarize()` can be input to graph

*#bar chart of number of events, all universities combined*

```
df_event %>% group_by(instnm, event_type) %>%  
  summarize(num_events=n()) %>%  
  ggplot(aes(x=event_type, y=num_events)) + # plot  
  ylab("Number of events") + xlab("Event type") +geom_col()
```

*#bar chart of number of events, separete chart for each university*

```
df_event %>% group_by(instnm, event_type) %>%  
  summarize(num_events=n()) %>%  
  ggplot(aes(x=event_type, y=num_events)) + # plot  
  ylab("Number of events") + xlab("Event type") + geom_col() +  
  coord_flip() + facet_wrap(~ instnm)
```

## Combining `summarize()` and `group_by`

**Task.** Count number of recruiting events by institution, event\_type, and whether event is in- or out-of-state (`var= event_inst` )

► Note: in `group_by()` , the optional `drop` argument controls whether empty groups dropped. default is `drop = TRUE`

XIN - WHY ISN'T `.DROP = FALSE` [BELOW] KEEPING ROWS FOR EMPTY GROUPS?

```
df_event %>% group_by(instnm, event_type, event_inst) %>%  
  summarize(num_events=n())
```

```
df_event %>% group_by(instnm, event_type, event_inst, .drop = TRUE) %>%  
  summarize(num_events=n())
```

```
df_event %>% group_by(instnm, event_type, event_inst, .drop = FALSE) %>%  
  summarize(num_events=n())
```

Make a graph, showing in/out state as fill color of bar

```
df_event %>% group_by(instnm, event_type, event_inst) %>%  
  summarize(num_events=n()) %>%  
  ggplot(aes(x=event_type, y=num_events, fill = event_inst)) + # plot  
  ylab("Number of events") + xlab("Event type") + geom_col() +  
  coord_flip() + facet_wrap(~ instnm)
```

## Combining `summarize()` and `group_by`

### Task

- By university, event type, event\_inst count the number of events and calculate the avg. pct white in the zip-code

```
df_event %>% group_by(instnm, event_type, event_inst) %>%  
  summarize(num_events=n(),  
            mean_pct_white=mean(pct_white_zip, na.rm = TRUE)  
  )
```

*#investigate object you created*

```
df_event %>% group_by(instnm, event_type, event_inst) %>%  
  summarize(num_events=n(),  
            mean_pct_white=mean(pct_white_zip, na.rm = TRUE)  
  ) %>% glimpse()
```



## Combining `summarize()` and `group_by`

Recruiting events by UC Berkeley

```
df_event %>% filter(univ_id == 110635) %>%  
  group_by(event_type) %>% summarize(num_events=n())
```

Let's create a dataset of recruiting events at UC Berkeley

```
event_berk <- df_event %>% filter(univ_id == 110635)  
  
event_berk %>% count(event_type)
```

summarize() and Counts

## `summarize()` : Counts

The count function `n()` takes no arguments and returns the size of the current group

```
event_berk %>% group_by(event_type, event_inst) %>%  
  summarize(num_events=n())
```

Because counts are so important, `dplyr` package includes separate `count()` function that can be called outside `summarize()` function

```
event_berk %>% group_by(event_type, event_inst) %>% count()
```

```
event_berk %>% group_by(event_type, event_inst) %>%  
  summarize(num_events=n()) # same same
```

## `summarize()` : count with logical vectors and `sum()`

Logical vectors have values `TRUE` and `FALSE` .

► When used with numeric functions, `TRUE` converted to 1 and `FALSE` to 0.

`sum()` is a numeric function that returns the sum of values

```
sum(c(5,10))
```

```
sum(c(TRUE,TRUE,FALSE,FALSE))
```

`is.na()` returns `TRUE` if value is `NA` and otherwise returns `FALSE`

```
is.na(c(5,NA,4,NA))
```

```
#> [1] FALSE TRUE FALSE TRUE
```

```
sum(is.na(c(5,NA,4,NA,5)))
```

```
#> [1] 2
```

```
sum(!is.na(c(5,NA,4,NA,5)))
```

```
#> [1] 3
```

Application: How many missing/non-missing obs in variable [**very important**]

```
event_berk %>% group_by(event_type) %>%
```

```
  summarize(
```

```
    n_events = n(),
```

```
    n_miss_inc = sum(is.na(med_inc)),
```

```
    n_nonmiss_inc = sum(!is.na(med_inc)),
```

```
    n_nonmiss_fr_lunch = sum(!is.na(fr_lunch))
```

```
)
```

## summarize() and count student exercise

Use one code chunk for this exercise. You could tackle this a step at a time and run the entire code chunk when you have answered all parts of this question. Create your own variable names.

1. Using the `event_berk` object, filter observations where `event_state` is VA and group by `event_type`.
  - 1.1 Using the `summarize` function to create a variable that represents the count for each `event_type`.
  - 1.2 Create a variable that represents the sum of missing obs for `med_inc`.
  - 1.3 create a variable that represents the sum of non-missing obs for `med_inc`.
  - 1.4 **Bonus:** Arrange variable you created representing the count of each `event_type` in descending order.

## summarize() and count student exercise SOLUTION

1. Using the `event_berk` object filter observations where `event_state` is VA and group by `event_type`.

- 1.1 Using the `summarize` function, create a variable that represents the count for each `event_type`.

- 1.2 Now get the sum of missing obs for `med_inc`.

- 1.3 Now get the sum of non-missing obs for `med_inc`.

```
event_berk %>%  
  filter(event_state == "VA") %>%  
  group_by(event_type) %>%  
  summarize(  
    n_events = n(),  
    n_miss_inc = sum(is.na(med_inc)),  
    n_nonmiss_inc = sum(!is.na(med_inc))) %>%  
  arrange(desc(n_events))  
  
#> # A tibble: 3 x 4  
#>   event_type n_events n_miss_inc n_nonmiss_inc  
#>   <chr>      <int>      <int>      <int>  
#> 1 public hs      20          0          20  
#> 2 private hs     13          0          13  
#> 3 other         3          0          3
```

summarize() and means

## summarize() : means

The `mean()` function within `summarize()` calculates means, separately for each group

```
event_berk %>% group_by(event_inst, event_type) %>% summarize(  
  n_events=n(),  
  mean_inc=mean(med_inc, na.rm = TRUE),  
  mean_pct_white=mean(pct_white_zip, na.rm = TRUE))
```

```
#> # A tibble: 10 x 5
```

```
#>   event_inst event_type  n_events mean_inc mean_pct_white  
#>   <chr>      <chr>      <int>    <dbl>      <dbl>  
#> 1 In-State  2yr college    111  78486.      40.1  
#> 2 In-State  4yr college     14 131691.      58.0  
#> 3 In-State  other           49  75040.      37.6  
#> 4 In-State  private hs       35  95229.      48.4  
#> 5 In-State  public hs      259  87097.      39.6  
#> 6 Out-State 2yr college      1 153070.      89.7  
#> 7 Out-State 4yr college      4  76913.      65.8  
#> 8 Out-State other          89  69004.      56.5  
#> 9 Out-State private hs    134  87654.      64.3  
#> 10 Out-State public hs    183 103603.      62.0
```



## `summarize()` : means and `na.rm` argument

Default behavior of “aggregation functions” (e.g., `summarize()` )

- ▶ if *input* has any missing values ( `NA` ), then output will be missing.

Many functions have argument `na.rm` (means “remove `NAs` ”)

- ▶ `na.rm = FALSE` [the default for `mean()` ]
  - ▶ Do not remove missing values from input before calculating
  - ▶ Therefore, missing values in input will cause output to be missing
- ▶ `na.rm = TRUE`
  - ▶ Remove missing values from input before calculating
  - ▶ Therefore, missing values in input will not cause output to be missing

*#na.rm = FALSE; the default setting*

```
event_berk %>% group_by(event_inst, event_type) %>% summarize(  
  n_events=n(),  
  n_miss_inc = sum(is.na(med_inc)),  
  mean_inc=mean(med_inc, na.rm = FALSE),  
  n_miss_frlunch = sum(is.na(fr_lunch)),  
  mean_fr_lunch=mean(fr_lunch, na.rm = FALSE))
```

*#na.rm = TRUE*

```
event_berk %>% group_by(event_inst, event_type) %>% summarize(  
  n_events=n(),  
  n_miss_inc = sum(is.na(med_inc)),  
  mean_inc=mean(med_inc, na.rm = TRUE),  
  n_miss_frlunch = sum(is.na(fr_lunch)),  
  mean_fr_lunch=mean(fr_lunch, na.rm = TRUE))
```

# Student exercise

1. Using the `event_berk` object, group by `instnm`, `event_inst`, & `event_type`.
  - 1.1 Create vars for number non\_missing for these racial/ethnic groups ( `pct_white_zip`, `pct_black_zip`, `pct_asian_zip`, `pct_hispanic_zip`, `pct_amerindian_zip`, `pct_nativehawaii_zip` )
  - 1.2 Create vars for mean percent for each racial/ethnic group

## Student exercise solutions

```
event_berk %>% group_by(instnm, event_inst, event_type) %>%
  summarize(
    n_events=n(),
    n_miss_white = sum(!is.na(pct_white_zip)),
    mean_white = mean(pct_white_zip, na.rm = TRUE),
    n_miss_black = sum(!is.na(pct_black_zip)),
    mean_black = mean(pct_black_zip, na.rm = TRUE),
    n_miss_asian = sum(!is.na(pct_asian_zip)),
    mean_asian = mean(pct_asian_zip, na.rm = TRUE),
    n_miss_lat = sum(!is.na(pct_hispanic_zip)),
    mean_lat = mean(pct_hispanic_zip, na.rm = TRUE),
    n_miss_na = sum(!is.na(pct_amerindian_zip)),
    mean_na = mean(pct_amerindian_zip, na.rm = TRUE),
    n_miss_nh = sum(!is.na(pct_nativehawaii_zip)),
    mean_nh = mean(pct_nativehawaii_zip, na.rm = TRUE)) %>%
  head(6)
```

```
#> # A tibble: 6 x 16
```

```
#>   instnm event_inst event_type n_events n_miss_white mean_white
#>   <chr>   <chr>      <chr>      <int>      <int>      <dbl>
#> 1 UC Be~ In-State 2yr colle~    111        106        40.1
#> 2 UC Be~ In-State 4yr colle~     14         12        58.0
#> 3 UC Be~ In-State other          49         48        37.6
#> 4 UC Be~ In-State private hs     35         35        48.4
#> 5 UC Be~ In-State public hs    259        258        39.6
#> 6 UC Be~ Out-State 2yr colle~      1          1        89.7
#> # ... with 10 more variables: n_miss_black <int>, mean_black <dbl>,
#> #   n_miss_asian <int>, mean_asian <dbl>, n_miss_lat <int>,
```

`summarize()` and logical vectors, part II

## summarize() : counts with logical vectors, part II

Logical vectors (e.g., `is.na()`) useful for counting obs that satisfy some condition

```
is.na(c(5,NA,4,NA))  
#> [1] FALSE TRUE FALSE TRUE  
typeof(is.na(c(5,NA,4,NA)))  
#> [1] "logical"  
sum(is.na(c(5,NA,4,NA)))  
#> [1] 2
```

**Task:** Using object `event_berk`, calculate the following measures for each combination of `event_type` and `event_inst`:

- ▶ count of number of rows for each group
- ▶ count of rows non-missing for both `pct_black_zip` and `pct_hispanic_zip`
- ▶ count of number of visits to communities where the `sum` of Black and Latinx people comprise more than 50% of the total population

```
event_berk %>% group_by (event_inst, event_type) %>%  
  summarize(  
    n_events=n(),  
    n_nonmiss_latbl = sum(!is.na(pct_black_zip) & !is.na(pct_hispanic_zip)),  
    n_majority_latbl= sum(pct_black_zip+ pct_hispanic_zip>50, na.rm = TRUE)  
  )
```

## `summarize()` : logical vectors to count *proportions*

Syntax: `group_by(vars) %>% summarize(prop = mean(TRUE/FALSE conditon))`

**Task:** separately for in-state/out-of-state, what proportion of visits to public high schools are to communities with median income greater than \$100,000?

Steps:

1. Filter public HS visits
2. group by in-state vs. out-of-state
3. Create measure

```
event_berk %>% filter(event_type == "public hs") %>% # filter public hs visits
  group_by (event_inst) %>% # group by in-state vs. out-of-state
  summarize(
    n_events=n(), # number of events by group
    n_nonmiss_inc = sum(!is.na(med_inc)), # w/ nonmissings values median inc,
    p_incgt100k = mean(med_inc>100000, na.rm=TRUE)) # proportion visits to $100k
#> # A tibble: 2 x 4
#>   event_inst n_events n_nonmiss_inc p_incgt100k
#>   <chr>      <int>      <int>      <dbl>
#> 1 In-State    259        256      0.273
#> 2 Out-State   183        183      0.519
```

`summarize()` : logical vectors to count *proportions*

What if we forgot to put `na.rm=TRUE` in the above task?

**Task:** separately for in-state/out-of-state, what proportion of visits to public high schools are to communities with median income greater than \$100,000?

```
event_berk %>% filter(event_type == "public hs") %>% # filter public hs visits
  group_by (event_inst) %>% # group by in-state vs. out-of-state
  summarize(
    n_events=n(), # number of events by group
    n_nonmiss_inc = sum(!is.na(med_inc)), # w/ nonmissings values median inc,
    p_incgt100k = mean(med_inc>100000, , na.rm=TRUE)) # proportion visits to $100k
#> # A tibble: 2 x 4
#>   event_inst n_events n_nonmiss_inc p_incgt100k
#>   <chr>      <int>      <int>      <dbl>
#> 1 In-State      259        256        0.273
#> 2 Out-State     183        183        0.519
```

## `summarize()` : Other “helper” functions

Lots of other functions we can use within `summarize()`

Common functions to use with `summarize()` :

Function	Description
<code>n</code>	count
<code>n_distinct</code>	count unique values
<code>mean</code>	mean
<code>median</code>	median
<code>max</code>	largest value
<code>min</code>	smallest value
<code>sd</code>	standard deviation
<code>sum</code>	sum of values
<code>first</code>	first value
<code>last</code>	last value
<code>nth</code>	nth value
<code>any</code>	condition true for at least one value

*Note: These functions can also be used on their own or with `mutate()`*



## summarize() : Other functions

Maximum value in a group

```
max(c(10,50,8))
```

```
#> [1] 50
```

**Task:** For each combination of in-state/out-of-state and event type, what is the maximum value of `med_inc` ?

```
event_berk %>% group_by(event_type, event_inst) %>%
```

```
  summarize(max_inc = max(med_inc)) # oops, we forgot to remove NAs!
```

```
#> # A tibble: 10 x 3
```

```
#>   event_type event_inst max_inc
```

```
#>   <chr>      <chr>      <dbl>
```

```
#> 1 2yr college In-State      NA
```

```
#> 2 2yr college Out-State 153070.
```

```
#> 3 4yr college In-State      NA
```

```
#> 4 4yr college Out-State      NA
```

```
#> 5 other      In-State      NA
```

```
#> 6 other      Out-State      NA
```

```
#> 7 private hs In-State 250001
```

```
#> 8 private hs Out-State      NA
```

```
#> 9 public hs  In-State      NA
```

```
#> 10 public hs Out-State 223556.
```

```
event_berk %>% group_by(event_type, event_inst) %>%
```

```
  summarize(max_inc = max(med_inc, na.rm = TRUE))
```

```
#> # A tibble: 10 x 3
```

```
#>   event_type event_inst max_inc
```

## summarize() : Other functions

Isolate first/last/nth observation in a group

```
x <- c(10,15,20,25,30)
first(x)
last(x)
nth(x,1)
nth(x,3)
nth(x,10)
```

**Task:** after sorting object `event_berk` by `event_type` and `event_datetime_start`, what is the value of `event_date` for:

- ▶ first event for each event type?
- ▶ the last event for each event type?
- ▶ the 50th event for each event type?

```
event_berk %>% arrange(event_type, event_datetime_start) %>%
  group_by(event_type) %>%
  summarize(
    n_events = n(),
    date_first= first(event_date),
    date_last= last(event_date),
    date_50th= nth(event_date, 50)
  )
```

# Student exercise

Identify value of `event_date` for the *nth* event in each by group

## Specific task:

- ▶ arrange (i.e., sort) by `event_type` and `event_datetime_start` , then group by `event_type` , and then identify the value of `event_date` for:
  - ▶ the first event in each by group ( `event_type` )
  - ▶ the second event in each by group
  - ▶ the third event in each by group
  - ▶ the fourth event in each by group
  - ▶ the fifth event in each by group

## Student exercise solution

```
event_berk %>% arrange(event_type, event_datetime_start) %>%
  group_by(event_type) %>%
  summarize(
    n_events = n(),
    date_1st= first(event_date),
    date_2nd= nth(event_date,2),
    date_3rd= nth(event_date,3),
    date_4th= nth(event_date,4),
    date_5th= nth(event_date,5))
#> # A tibble: 5 x 7
#>   event_type n_events date_1st   date_2nd   date_3rd   date_4th
#>   <chr>      <int> <date>     <date>     <date>     <date>
#> 1 2yr colle~    112 2017-04-25 2017-09-05 2017-09-05 2017-09-06
#> 2 4yr colle~     18 2017-04-30 2017-05-01 2017-05-06 2017-09-13
#> 3 other       138 2017-04-11 2017-04-23 2017-04-25 2017-04-29
#> 4 private hs   169 2017-04-23 2017-04-24 2017-04-29 2017-04-30
#> 5 public hs    442 2017-04-14 2017-04-24 2017-04-26 2017-04-27
#> # ... with 1 more variable: date_5th <date>
```

Summarize multiple columns

## What are “scoped” variants of a function?

“Scoped” variants of a function apply the function to a selection of variables.

Three kinds of scoped variants exist:

1. Verbs (i.e., functions) suffixed with `_all()` apply an operation on all variables.
  - ▶ e.g.: `summarize_all()`, `mutate_all()`
2. Verbs suffixed with `_at()` (e.g., `summarize_at()`) apply an operation on a subset of variables specified with quoting function `vars()`.
  - ▶ This quoting function accepts helpers functions like `starts_with()`
3. Verbs suffixed with `_if()` apply an operation on the subset of variables for which a predicate function returns TRUE.

Arguments of scoped variants

- ▶ `.tbl` A tbl object (data frame)
- ▶ `.funs` specifies which function(s) to perform (e.g., calculate mean)
  - ▶ Argument values: A function `fun`; a quosure style lambda `~ fun(.)`; or a list of either form (e.g., `'list(mean,min,max)`).
- ▶ `.vars` which variables to apply function to:
  - ▶ argument values: A list of columns generated by `vars()`, a character vector of column names, a numeric vector of column positions, or `NULL`.
- ▶ `.predicate` A predicate function to be applied to the columns or a logical vector. The variables for which `.predicate` is or returns `TRUE` are selected.
- ▶ `...` Additional arguments for function calls in `.funs`, evaluated once w/ tidy dots support

# What are “scoped” variants of a function?

Why/when use “scoped” variants of a function

- ▶ When you want to perform an operation on multiple variables without naming each individual variable

“verbs” (i.e., functions) from the `dplyr` package that have scoped variants `_all()`, `_at()`, and `_if()`

- ▶ `mutate()` and `transmute()` [see `?mutate_all`]
- ▶ `summarize()` [see `?summarize_all`]
- ▶ `filter()`
- ▶ `group_by()`
- ▶ `rename()` and `select()`
- ▶ `arrange()`

## Scoped variants of summarize()

**Description.** The “scoped variants” of `summarize()` apply `summarize()` to multiple variables. Three variants:

- ▶ `summarize_all()` affects every variable
- ▶ `summarize_at()` affects variables selected with a character vector or `vars()`
- ▶ `summarize_if()` affects variables selected with a predicate function

### Syntax

- ▶ `summarize_all(.tbl, .funcs, ...)`
- ▶ `summarize_at(.tbl, .vars, .funcs, ...)`
- ▶ `summarize_if(.tbl, .predicate, .funcs, ...)`

### Arguments

- ▶ `.tbl` A `tbl` object (data frame)
- ▶ `.funcs` specifies which function(s) to perform (e.g., calculate mean)
  - ▶ Argument values: A function `fun`; a quosure style lambda `~ fun(.)`; or a list of either form (e.g., `'list(mean,min,max)'`).
- ▶ `.vars` which variables to apply function to:
  - ▶ argument values: A list of columns generated by `vars()`, a character vector of column names, a numeric vector of column positions, or `NULL`.
- ▶ `.predicate` A predicate function to be applied to the columns or a logical vector. The variables for which `.predicate` is or returns `TRUE` are selected.
- ▶ `...` Additional arguments for the function calls in `.funcs`.
  - ▶ These are evaluated only once, with tidy dots support.



## summarize\_all() affects every variable

Syntax: `summarize_all(.tbl, .funcs, ...)`

- ▶ `.tbl` A tbl object (data frame)
- ▶ `.funcs` specifies which function(s) to perform. Argument values:
  - ▶ A function `fun` ; a quosure style lambda `~ fun(.)`; a list (e.g., `list(mean,min)`)
- ▶ `...` Additional arguments for function calls in `.funcs` . These are evaluated once

### Task:

- ▶ For U. Pittsburgh ( `univ_id = 215293` ) events at public and private high schools, calculate the **mean** value of `med_inc` and `pct_white_zip` for each combination of `event_type` and `event_inst`

```
df_event %>%  
  filter(univ_id == 215293, event_type %in% c("private hs", "public hs")) %>%  
  select(event_type, event_inst, med_inc, pct_white_zip) %>%  
  group_by(event_type, event_inst) %>%  
  summarize_all(.funcs = mean)
```

Try again, this time applying `na.rm = TRUE`

- ▶ this is an example of a `...` argument “for the function calls in `.funcs` .”

```
df_event %>%  
  filter(univ_id == 215293, event_type %in% c("private hs", "public hs")) %>%  
  select(event_type, event_inst, med_inc, pct_white_zip) %>%  
  group_by(event_type, event_inst) %>%  
  summarize_all(.funcs = mean, na.rm = TRUE)
```

## summarize\_all() affects every variable

Syntax: `summarize_all(.tbl, .funs, ...)`

- ▶ `.tbl` A tbl object (data frame)
- ▶ `.funs` specifies which function(s) to perform. Argument values:
  - ▶ A function `fun` ; a quosure style lambda `~ fun(.)`; a list (e.g., `list(mean,min)`)
- ▶ `...` Additional arguments for function calls in `.funs` . These are evaluated once

### Task:

- ▶ For U. Pittsburgh ( `univ_id = 215293` ) events at public and private high schools, calculate **mean** and **standard deviation** of `med_inc` and `pct_white_zip` for each combination of `event_type` and `event_inst`

```
df_event %>%  
  filter(univ_id == 215293, event_type %in% c("private hs", "public hs")) %>%  
  select(event_type, event_inst, med_inc, pct_white_zip) %>%  
  group_by(event_type, event_inst) %>%  
  summarize_all(.funs = list(mean, sd), na.rm = TRUE)
```

Use this syntax to control variable name suffixes:

- ▶ `.funs = list(var_name_suffix = function_name, ...)`

```
df_event %>%  
  filter(univ_id == 215293, event_type %in% c("private hs", "public hs")) %>%  
  select(event_type, event_inst, med_inc, pct_white_zip) %>%  
  group_by(event_type, event_inst) %>%  
  summarize_all(.funs = list(avg = mean, std = sd), na.rm = TRUE)
```

## summarize\_all() affects every variable

### Task:

- ▶ Same task as before, but now calculate, **mean**, **standard deviation**, **\_\_min\_\_**, and **max** of `med_inc` and `pct_white_zip` for each combination of `event_type` and `event_inst`

```
df_event %>%  
  filter(univ_id == 215293, event_type %in% c("private hs", "public hs")) %>%  
  select(event_type, event_inst, med_inc, pct_white_zip) %>%  
  group_by(event_type, event_inst) %>%  
  summarize_all(.funs = list(avg = mean, std = sd, low = min, high = max), na.rm = TRUE)
```

## summarize\_at() affects selected variables

Syntax: `summarize_at(.tbl, .vars, .funcs, ...)`

- ▶ `.tbl` A tbl object (data frame)
- ▶ `.vars` which variables to operate on. Argument values:
  - ▶ A list of columns generated by `vars()`, a character vector of column names, a numeric vector of column positions, or `NULL`.
- ▶ `.funcs` specifies which function(s) to perform. Argument values:
  - ▶ A function `fun` ; a quosure style lambda `~ fun(.)`; a list (e.g., `list(mean,min)`)
- ▶ `...` Additional arguments for function calls in `.funcs` . These are evaluated once

**Task:** For U. Pittsburgh events at public and private high schools, calculate **mean**, **min**, and **max** of variables `med_inc` and `event_date` for each combination of `event_type` and `event_inst`

```
df_event %>%  
  filter(univ_id == 215293, event_type %in% c("private hs", "public hs")) %>%  
  group_by(event_type, event_inst) %>%  
  summarize_at(.vars = vars(med_inc, event_date),  
               .funcs = list(avg = mean, low = min, high = max), na.rm = TRUE)
```

alternative:

```
df_event %>%  
  filter(univ_id == 215293, event_type %in% c("private hs", "public hs")) %>%  
  group_by(event_type, event_inst) %>%  
  summarize_at(.vars = c("med_inc", "event_date"),  
               .funcs = list(avg = mean, low = min, high = max), na.rm = TRUE)
```

## summarize if

The `_if()` variants apply a predicate function (a function that returns TRUE or FALSE) to determine the relevant subset of columns. Here we apply `mean()` to the numeric columns: `starwars %>% summarize_if(is.numeric, mean, na.rm = TRUE)`

Attach aggregate measures to your data frame

## Attach aggregate measures to your data frame

We can attach aggregate measures to a data frame by using `group_by` without `summarize()`

What do I mean by “attaching aggregate measures to a data frame”?

- ▶ Calculate measures at the `by_group` level, but attach them to original object rather than creating an object with one row for each `by_group`

**Task:** Using `event_berk` data frame, create (1) a measure of average income across all events and (2) a measure of average income for each event type

- ▶ resulting object should have same number of observations as `event_berk`

Steps:

1. create measure of avg. income across all events without using `group_by()` or `summarize()` and assign as (new) object
2. Using object from previous step, create measure of avg. income across by event type using `group_by()` without `summarize()` and assign as new object

## Attach aggregate measures to your data frame

**Task:** Using `event_berk` data frame, create (1) a measure of average income across all events and (2) a measure of average income for each event type

1. Create measure of average income across all events

```
event_berk_temp <- event_berk %>%  
  arrange(event_date) %>% # sort by event_date (optional)  
  select(event_date, event_type, med_inc) %>% # select vars to be retained (optional)  
  mutate(avg_inc = mean(med_inc, na.rm=TRUE)) # create avg. inc measure  
  
dim(event_berk_temp)  
event_berk_temp %>% head(5)
```

2. Create measure of average income by event type

```
event_berk_temp <- event_berk_temp %>%  
  group_by(event_type) %>% # grouping by event type  
  mutate(avg_inc_type = mean(med_inc, na.rm=TRUE)) # create avg. inc measure  
  
str(event_berk_temp)  
event_berk_temp %>% head(5)
```



## Attach aggregate measures to your data frame

**Task:** Using `event_berk_temp` from previous question, create a measure that identifies whether `med_inc` associated with the event is higher/lower than average income for all events of that type

Steps:

1. Create measure of average income for each event type [already done]
2. Create 0/1 indicator that identifies whether median income at event location is higher than average median income for events of that type

```
# average income at recruiting events across all universities
event_berk_tempv2 <- event_berk_temp %>%
  mutate(gt_avg_inc_type = med_inc > avg_inc_type) %>%
  select(-(avg_inc)) # drop avg_inc (optional)
event_berk_tempv2 # note how med_ic = NA are treated
```

Same as above, but this time create integer indicator rather than logical

```
event_berk_tempv2 <- event_berk_tempv2 %>%
  mutate(gt_avg_inc_type = as.integer(med_inc > avg_inc_type))
event_berk_tempv2 %>% head(4)
```

## Student exercise

Task: is `pct_white_zip` at a particular event higher or lower than the average `pct_white_zip` for that `event_type`?

- ▶ Note: all events attached to a particular `zip_code`
- ▶ `pct_white_zip`: pct of people in that `zip_code` who identify as white

Steps in task:

- ▶ Create measure of average pct white for each `event_type`
- ▶ Compare whether `pct_white_zip` is higher or lower than this average

## Student exercise solution

Task: is `pct_white_zip` at a particular event higher or lower than the average `pct_white_zip` for that `event_type` ?

```
event_berk_tempv3 <- event_berk %>%  
  arrange(event_date) %>% # sort by event_date (optional)  
  select(event_date, event_type, pct_white_zip) %>% #optional  
  group_by(event_type) %>% # grouping by event type  
  mutate(avg_pct_white = mean(pct_white_zip, na.rm=TRUE),  
         gt_avg_pctwhite_type = as.integer(pct_white_zip > avg_pct_white))  
event_berk_tempv3 %>% head(4)  
#> # A tibble: 4 x 5  
#>   event_date event_type pct_white_zip avg_pct_white gt_avg_pctwhite_type  
#>   <date>      <chr>          <dbl>         <dbl>          <int>  
#> 1 2017-04-11 other           37.2          49.7            0  
#> 2 2017-04-14 public hs       78.3          48.9            1  
#> 3 2017-04-23 private hs      84.7          61.0            1  
#> 4 2017-04-23 other          20.9          49.7            0
```